

Verification with Sharing and Aliasing

Gregory Malecha

Harvard University SEAS

March 5, 2010

Why Verify?

Observation

All large, complex systems have bugs.

- Hardware design – Intel floating point bug (\$300 million)
- Mars rover – Priority inversion
- Security – Internet viruses and worms
- Voting machines – Hacking voting machines (an election?)
- Safety – control systems for airplanes, power plants, space shuttle

A problem that isn't going away...

- Just waiting won't solve this problem...
 - A computer that runs twice as fast will just trigger twice as many bugs per second.

A problem that isn't going away...

- Just waiting won't solve this problem...
 - A computer that runs twice as fast will just trigger twice as many bugs per second.
- ...actually time is making it harder.
 - Multicore and multiprocessor means reasoning about concurrency.
 - Lax memory models make low-level reasoning more difficult.

A problem that isn't going away...

- Just waiting won't solve this problem...
 - A computer that runs twice as fast will just trigger twice as many bugs per second.
- ...actually time is making it harder.
 - Multicore and multiprocessor means reasoning about concurrency.
 - Lax memory models make low-level reasoning more difficult.
- And, we're trying to solve bigger problems than before...
 - Data integrity and security
 - Scientific simulation


Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

Outline


- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

Techniques for Reasoning About Code



- Basic trade-off between the amount of effort required and the expressivity of the properties.


Techniques for Reasoning About Code



Goal


- Strong guarantees about complex properties.
- Scalable and modular.

Techniques for Reasoning About Code




- We would love it if just looking at the code was here...

Techniques for Reasoning About Code



- We would love it if just looking at the code was here...
- But we all know it's more like here...


Techniques for Reasoning About Code



Testing

- **Pro** Direct and intuitive methodology.
- **Con** Large amount of code, very manual.


Techniques for Reasoning About Code



Type Systems

- **Pro** Fast, (can be) provably correct and compositional.
- **Con** Limited properties, “restricted programming”.


Techniques for Reasoning About Code



Model Checking

- **Pro** “Push-button” when it works and somewhat intuitive.
- **Con** Computationally expensive, can be difficult to set up.


Techniques for Reasoning About Code



Theorem Proving

- **Pro** Expressive and provably correct.
- **Con** Proving can be tedious, often requires an “expert”.

Techniques for Reasoning About Code



Theorem Proving

- **Pro** Expressive and provably correct.
- **Con** Proving can be tedious, often requires an “expert”.
- This is the focus of the talk.


The Myth of “Correctness”

- “Correct” is dependent on what the system should do.



The Myth of “Correctness”


- “Correct” is dependent on what the system should do.



- Errors can enter at the specification level.
 - Specification shouldn't talk about complex implementation details.
 - Should be easier to write and reason about.

The Myth of “Correctness”


- “Correct” is dependent on what the system should do.



- Errors can enter at the specification level.
 - Specification shouldn't talk about complex implementation details.
 - Should be easier to write and reason about.
- We can verify an implementation with respect to a specification.

The Myth of “Correctness”

- “Correct” is dependent on what the system should do.



- Errors can enter at the specification level.
 - Specification shouldn't talk about complex implementation details.
 - Should be easier to write and reason about.
- We can verify an implementation with respect to a specification.
- Compile the implementation in a certified way.

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

Building on Types

- How do you figure out what a function does?

```
int largest(int cnt, int* ary) {  
    ... /** implementation */ ...  
}
```

Building on Types

- How do you figure out what a function does?

```
int largest(int cnt, int* ary) {  
    ... /** implementation */ ...  
}
```

Building on Types

- How do you figure out what a function does?

```
/** largest(cnt, ary)  
**   returns the largest element in the first  
**   cnt elements of ary  
** Requires:  
** = 1 <= cnt <= length of ary  
**/  
int largest(int cnt, int* ary) {  
    ... /** implementation */ ...  
}
```


A Little More Expressive

- Annotation languages like *PREfix*/*PREfast* allow specifying properties like array bounds information.

```

/** largest(cnt, ary)
    **   returns the largest element in the first
    **   cnt elements of ary
    ** Assumes:
    ** = 1 <= cnt <= length of ary
    **/
int largest(int cnt, __in__ecount(cnt) int* ary) {
    ... /** implementation **/ ...
}

```

Specifications as Dependent Types

- Still aren't specifying everything...
 - Input: Empty arrays.
 - Output: The result is really the largest element.

```
largest(int cnt, int[cnt] ary, (0 < cnt) _pf) :  
  {x : int | maximal x ary}  
{ ... /* implementation */ ... }
```

- Types *depend* on run-time values.
 - Length of ary is cnt.
- Require proofs of preconditions & return proofs of correctness.
 - Proof that $0 < \text{cnt}$.
 - Returns pair of the result and a proof that the result is correct.

A Monkey-Wrench: Effects

- The previous code was basically functional.
- Most programs use imperative state and effects.

```
void sortInPlace(int cnt, int[] ary) {  
    ... /** Implementation */ ...  
}
```

- We need to state that the contents of ary changes.

A Standard Approach

- Can reason about effectful code using Hoare Logic.

$$\{P\} c \{r \Rightarrow Q\}$$

- P is the precondition.
 - c is the command to execute.
 - r is a binder for the return value.
 - Q is the postcondition which depends on r .
- When the state of the program is described by P , c can be run and, if c terminates with return value r , the state of the program will be described by Q .

Describing the World

Example Program

```
{ p1 ↦ 1 ∧ p2 ↦ 1 }  
*p1 = 3  
{ _ ⇒ p1 ↦ 3 ∧ p2 ↦ 1 }
```

- Can we prove this?

Describing the World

Example Program

```
{ p1 ↦ 1 ∧ p2 ↦ 1 }
*p1 = 3
{ _ ⇒ p1 ↦ 3 ∧ p2 ↦ ??? }
```

- Can we prove this? No.
 - What if p_1 is an alias of p_2 ?


Describing the World

Example Program

```
{ p1 ↦ 1 ∧ p2 ↦ 1 ∧ [p1 ≠ p2] }  
*p1 = 3  
{ _ ⇒ p1 ↦ 3 ∧ p2 ↦ 1 }
```


- Can we prove this? No.
 - What if p_1 is an alias of p_2 ?
 - We need a side condition for every pair of pointers.
 - Can't encode abstraction easily.

Separation Logic (Reynolds '02)




- We want to be able to reason about two structures independently.

Separation Logic (Reynolds '02)



- We want to be able to reason about two structures independently.
- Encode the disjointness condition in the $*$.
 - Easy to write the common case when pointers don't alias.


Separation Logic (Reynolds '02)



- We want to be able to reason about two structures independently.
- Encode the disjointness condition in the $*$.
 - Easy to write the common case when pointers don't alias.
- Called the “Frame Rule”
 - Allows us to temporarily “forget” about the list, reason about the tree, and then remember the list.

$$\frac{}{\{tree * list\}c\{r \Rightarrow tree' * list\}} \text{Frame}$$


Separation Logic (Reynolds '02)



- We want to be able to reason about two structures independently.
- Encode the disjointness condition in the $*$.
 - Easy to write the common case when pointers don't alias.
- Called the “Frame Rule”
 - Allows us to temporarily “forget” about the list, reason about the tree, and then remember the list.

$$\frac{\{tree\}c\{r \Rightarrow tree'\}}{\{tree * list\}c\{r \Rightarrow tree' * list\}} \text{ Frame}$$

Separation Logic (Reynolds '02)



- We want to be able to reason about two structures independently.
- Encode the disjointness condition in the $*$.
 - Easy to write the common case when pointers don't alias.
- Called the “Frame Rule”
 - Allows us to temporarily “forget” about the list, reason about the tree, and then remember the list.

$$\frac{\{P\}c\{r \Rightarrow Q\}}{\{P * R\}c\{r \Rightarrow Q * R\}} \text{ Frame}$$

Hoare Type Theory: Specifying Effects in Types

- Embed Hoare logic into the *types* of terms.
- Hoare triples are represented by the `Cmd` type.

$$\{P\} c \{r \Rightarrow Q\} \equiv c : \text{Cmd } P (r \Rightarrow Q)$$

Pointer Operations

```
Write p v : Cmd (∃ w, p ↦ w)
              (_ ⇒ p ↦ v)
```

sortInPlace in HTT

```
sortInPlace(int cnt, int[cnt] ary, #list int# m)
  : Cmd (array ary m)
      (_  $\Rightarrow$  array ary (sort m))
{ ... /** implementation **/ ... }
```

- m is computationally irrelevant, i.e. compile-time only.
 - Used only to simplify reasoning.
- array a 1 is an abstraction predicate that states the contents of the array (a) are the same as the contents of the list (l).

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - ① How do we describe the value of the list?

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - ① How do we describe the value of the list?
Relate to an irrelevant list

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - 1 How do we describe the value of the list?
Relate to an irrelevant list
 - 2 How do we describe the heap that contains a particular list?

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - 1 How do we describe the value of the list?
Relate to an irrelevant list
 - 2 How do we describe the heap that contains a particular list?
Specify a representation predicate

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```

- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - 1 How do we describe the value of the list?
Relate to an irrelevant list
 - 2 How do we describe the heap that contains a particular list?
Specify a representation predicate
 - 3 What does each function do? What does “correct” mean?

Example: C-style Linked Lists

```
interface IntList {  
    Integer get(int index);  
    void insert(int index, int value);  
    ...  
}
```


- Specifies an abstract type `IntList` with two methods.
- To reason about correctness, we need specifications.
 - ① How do we describe the value of the list?
Relate to an irrelevant list
 - ② How do we describe the heap that contains a particular list?
Specify a representation predicate
 - ③ What does each function do? What does “correct” mean?
Give a Hoare-logic specification

An Elaborated Interface

```
Interface IntList H {  
  llist : H  $\rightarrow$  list int  $\rightarrow$  hprop ;  
  
  get (H h, int index, #list int# m)  
    : Cmd (llist h m)  
      (r  $\Rightarrow$  llist h m * [r = nth m index]) ;  
  insert (H h, int index, int val, #list int# m)  
    : Cmd (llist h m)  
      (_  $\Rightarrow$  llist h (spec_insert m index val)) ;  
  
  /** ... **/  
}
```

- H is the type of handles to lists.
- llist is the representation predicate.

Implementation: The Representation Predicate



- Describe the heap computationally using a functional model.


Implementation: The Representation Predicate

pStart ————— pEnd

- Describe the heap computationally using a functional model.

`llseg pStart pEnd nil` \iff `[pStart = pEnd]`

Implementation: The Representation Predicate




- Describe the heap computationally using a functional model.

```
Record llNode := mkNode { val : int ; next : optr }
```

```
llseg pStart pEnd nil  $\iff$  [pStart = pEnd]
```

```
llseg (Ptr p) pEnd (a :: b)  $\iff$   $\exists$  nx : optr ,  
p  $\mapsto$  mkNode a nx * llseg nx pEnd b
```

Implementation: The Representation Predicate



- Describe the heap computationally using a functional model.

```
Record llNode := mkNode { val : int ; next : optr }
```

```
llseg pStart pEnd nil  $\iff$  [pStart = pEnd]
```

```
llseg (Ptr p) pEnd (a :: b)  $\iff$   $\exists$  nx : optr ,
    p  $\mapsto$  mkNode a nx * llseg nx pEnd b
```

```
tlst  $\equiv$  ptr
```

```
llist h m  $\iff$   $\exists$  st : optr , h  $\mapsto$  st * llseg st Null m
```

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

An Interface for Iterators

- Iterators and collections go hand-in-hand.

```
Interface ListIterable titr {  
  iter : titr → list int → nat → hprop ;
```


An Interface for Iterators

- Iterators and collections go hand-in-hand.

```
Interface ListIterable titr {  
  iter : titr → list int → nat → hprop ;  
  next (titr t, #list int# m, #nat# index)  
    : Cmd (iter t m idx)  
      (res ⇒ iter t m (nextIndex index (length m)) *  
        [res = nth m index])  
}
```

- T is the type of values being iterated over.
- `titr` is the type of the iterator handle.
- Representation predicate (`iter`) and `next` command.


Implementing Iterators over Lists



`titr \equiv ptr`

`iter (t : titr) (ls : list int) (idx : nat) \iff
 \exists st : optr, \exists cur : optr, $t \mapsto (st, cur)$`

Implementing Iterators over Lists




$\text{titr} \equiv \text{ptr}$

```

iter (t : titr) (ls : list int) (idx : nat)  $\iff$ 
   $\exists$  st : optr,  $\exists$  cur : optr,  $t \mapsto (\text{st}, \text{cur})$  *
  llseg st cur (firstn idx ls)

```

Implementing Iterators over Lists



$\text{titr} \equiv \text{ptr}$


```

iter (t : titr) (ls : list int) (idx : nat)  $\iff$ 
   $\exists$  st : optr,  $\exists$  cur : optr,  $t \mapsto (\text{st}, \text{cur})$  *
  llseg st cur (firstn idx ls) *
  llseg cur Null (skipn idx ls)

```


The Sharing Problem

- Requires access to the same memory as the underlying list.
 - Creating an iterator transfers ownership of memory from the list to iterator.
 - Can't have multiple iterators.




The Sharing Problem

- Requires access to the same memory as the underlying list.
 - Creating an iterator transfers ownership of memory from the list to iterator.
 - Can't have multiple iterators.



The Sharing Problem

- Requires access to the same memory as the underlying list.
 - Creating an iterator transfers ownership of memory from the list to iterator.
 - Can't have multiple iterators.




The Sharing Problem: Specifications

- Computations on iterators can't be called with the same underlying list.

```
zip(titr i1, titr i2, #list int# l1, #list int# l2) :  
  Cmd (iter i1 l1 0 * iter i2 l2 0 * [length l1 = length l2])  
    (res ⇒  
      iter i1 l1 (length l1) * iter i2 l2 (length l2) *  
      llist res (fzip l1 l2))
```

A Real Sharing Problem

- Who “owns” the list turns out to be a real problem.




- Consider the following program:

```
Iterator<Integer> itr = lst.iterator();
```

A Real Sharing Problem

- Who “owns” the list turns out to be a real problem.




- Consider the following program:

```
Iterator<Integer> itr = lst.iterator();  
itr.next();
```


A Real Sharing Problem

- Who “owns” the list turns out to be a real problem.




- Consider the following program:

```
Iterator<Integer> itr = lst.iterator();  
itr.next();  
lst.remove(1);
```

A Real Sharing Problem

- Who “owns” the list turns out to be a real problem.




- Consider the following program:

```
Iterator<Integer> itr = lst.iterator();  
itr.next();  
lst.remove(1);  
itr.next();
```

A Real Sharing Problem

- Who “owns” the list turns out to be a real problem.




- Consider the following program:

```
Iterator<Integer> itr = lst.iterator();  
itr.next();  
lst.remove(1);  
lst.insert(1, 'B');  
itr.next();
```

- Source of Java's ConcurrentModificationException.

Sharing with Fractional Permissions (Boyland '03)

- Parametrize points-to by a fractional ownership.
 - $p \mapsto^q v$, q is the fraction.
- Ownership determines your capabilities:
 - Full permissions allows everything: read, write, free.
 - Partial permissions only allows reading.
 - Permissions can be split and joined.



A Fractional Iterator

- Describe the iterator as owning a fraction of the whole list.

```
(** Representation predicate **)
liter (owner : tlst) (q : Fp)
  (t : titr) (ls : list int) (idx : nat)  $\iff$ 
   $\exists$  st : optr,  $\exists$  cur : optr,
  owner  $\xrightarrow{q}$  st * t  $\xrightarrow{1}$  cur *
  llseg st cur (firstn idx ls) q *
  llseg cur Null (skipn idx ls) q
```

- q is the fraction of the list that is owned.
- Allows multiple iterators over the same list.

Exposing Fractions

- Need to prove that lists can be split...
 - $q \mid \# \mid q'$ states that q and q' are compatible, i.e. sum to less than or equal to 1.

Lemma `llist_perm_split` : $\forall q q' t \text{ ls},$
 $q \mid \# \mid q' \rightarrow$
 $\text{llist } q + q' t \text{ ls} \implies \text{llist } q t \text{ ls} * \text{llist } q' t \text{ ls}$

Exposing Fractions

- Need to prove that lists can be split...
 - $q \mid\# \mid q'$ states that q and q' are compatible, i.e. sum to less than or equal to 1.

Lemma `llist_perm_split` : $\forall q q' t \text{ ls},$
 $q \mid\# \mid q' \rightarrow$
 $\text{llist } q + q' t \text{ ls} \implies \text{llist } q t \text{ ls} * \text{llist } q' t \text{ ls}$

- ...and joined together.

Lemma `llist_perm_join` : $\forall q q' t \text{ ls},$
 $q \mid\# \mid q' \rightarrow$
 $\text{llist } q t \text{ ls} * \text{llist } q' t \text{ ls} \implies \text{llist } q + q' t \text{ ls}$

Recap: Fractional Iterators


- **Original Problem** Couldn't have multiple views of the same list.
 - Either a list or an iterator, not both.
 - Only 1 iterator at a time.
- **Solution** Fractional permissions allow sharing.
 - Lift fractional permissions to the level of abstract data types.
 - Only slight modifications to incorporate fractions.
 - Prove two simple lemmas about splitting and joining.
 - Able to pass-out read-only permissions, finer granularity permissions.

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions


Specifications with Aliasing

- Aliasing presents a unique problem for separation logic.
 - Lists are easy...




Specifications with Aliasing

- Aliasing presents a unique problem for separation logic.
 - Lists are easy...
 - Trees are easy...




Specifications with Aliasing

- Aliasing presents a unique problem for separation logic.
 - Lists are easy...
 - Trees are easy...
 - Trees with lists are not easy ...




Specifications with Aliasing

- Aliasing presents a unique problem for separation logic.
 - Lists are easy...
 - Trees are easy...
 - Trees with lists are not easy **because of aliasing**...



B+ Trees




- B+ trees are n -ary trees where the leaves are connected by a linked list.
 - Support fast lookup and in-order iteration.
 - Commonly used for database indices. (Malecha '10)
- Previous formalizations exist, but neither is mechanically verified:
 - Classical conjunction, $(\text{list} * \text{any}) \wedge (\text{tree})$. (Bornat '04)
 - B+ tree language. (Sexton '08)
- Both of these approaches seem difficult to automate.

Difficulties of the Invariant

- Have to encode pointer aliasing explicitly.


Difficulties of the Invariant

- Have to encode pointer aliasing explicitly.
- Many different B+ trees can describe the same finite map.




Difficulties of the Invariant

- Have to encode pointer aliasing explicitly.
- Many different B+ trees can describe the same finite map.



Difficulties of the Invariant

- Have to encode pointer aliasing explicitly.
- Many different B+ trees can describe the same finite map.



- Other properties that we won't focus on.
 - Enforce the tree balancedness.
 - Enforce the ordering of keys.
 - Invariants on the size of branches and leaves.

Defining a Representation Model

- A standard, functionally n -ary tree is enough for the trunk.

```
tree = Branch (list tree) | Leaf (list value)
```

Defining a Representation Model

- A standard, functionally n -ary tree is enough for the trunk.
`tree = Branch (list tree) | Leaf (list value)`
- This stores the structure, but the aliasing is still difficult.
 - We need to give equations on pointers, in the representation.

Defining a Representation Model

- A standard, functionally n -ary tree is enough for the trunk.

```
tree = Branch (list tree) | Leaf (list value)
```

- This stores the structure, but the aliasing is still difficult.
 - We need to give equations on pointers, in the representation.

- **Solution:** Elaborate the functional tree with the pointers.

```
ptree = Branch ptr * (list tree)  
       | Leaf ptr * (list value)
```

- Enforce that the pointer stored in each node is the pointer that points to the node.
 - Quantifies all pointers simultaneously.
 - Makes it easy to state aliasing constraints.

Representation Invariant

- Existentially quantify an irrelevant model (tr) of the tree which contains the pointers.
 - Avoids existentials in the representation invariant, simplifies automation.
 - Makes the heap predicate ($repTree$) very computational.

```
rep (p : BptMap) (m : Model)  $\iff$   
   $\exists$  pRoot : ptr,  $\exists$  tr : ptree,  
  p  $\mapsto$  (pRoot, #tr#) *  
  repTree pRoot Null tr
```

Representation Invariant

- Existentially quantify an irrelevant model (tr) of the tree which contains the pointers.
 - Avoids existentials in the representation invariant, simplifies automation.
 - Makes the heap predicate (repTree) very computational.
 - Connect the logical model (m) to the physical model (tr).


```
rep (p : BptMap) (m : Model)  $\iff$   
   $\exists$  pRoot : ptr,  $\exists$  tr : ptree,  
  p  $\mapsto$  (pRoot, #tr#) *  
  repTree pRoot Null tr *  
  [m = as_map tr]
```

Representation Invariant

- Existentially quantify an irrelevant model (tr) of the tree which contains the pointers.
 - Avoids existentials in the representation invariant, simplifies automation.
 - Makes the heap predicate (repTree) very computational.
 - Connect the logical model (m) to the physical model (tr).
 - Consolidate pure facts about the model in inv .


```
rep (p : BptMap) (m : Model)  $\iff$   
   $\exists$  pRoot : ptr,  $\exists$  tr : ptrtree,  
  p  $\mapsto$  (pRoot, #tr#) *  
  repTree pRoot Null tr *  
  [m = as_map tr] *  
  [inv tr MinK MaxK]
```


Implementation: insert and lookup



- Most operations act on the tree.
 - Efficient lookup ($O(\lg n)$).
 - Efficient insert ($O(\lg n)$).
- Implementation follows recursive structure of the tree
 - Simple recursion invariant.
 - Relatively simple to verify.
 - The complexities come from the width of the branches.

Implementation: Iteration



- Can switch between views by proving and applying a lemma:

```

Lemma repTree_repTrunkLeaves : ∀ (h : nat)
  (p : ptr) (last : optr) (m : ptree h),
  repTree p last m
  ⇔
  repTrunk p last m *
  repLeaves (Ptr (firstPtr m)) (leaves m) last.

```

Recap: B+ trees

- **Original Problem** Aliasing at the leaves and relational heap predicate makes describing the heap difficult.
 - Existing approaches seem cumbersome to verify.
- **Solution** Factor out the relation by quantifying an irrelevant model.
 - Including the pointers in the model makes them easy to access.
 - Simple, computational heap predicate.
 - Support multiple views by proving an equivalence of formulae.
 - Avoid unnecessary guessing during proof search.
 - Use irrelevance to avoid run-time overhead.

Outline

- 1 Techniques for Gaining Confidence
- 2 Software Verification with Types
 - Modularity and Abstraction
- 3 My Work: Addressing Sharing and Aliasing
 - Sharing: Iterators
 - Aliasing: B+ Trees
- 4 Conclusions

The Take-away

- Higher-order abstraction simplifies specifications and proofs.
- Fractional permissions are necessary even for sequential code.
- Separation logic makes trees much easier than DAGs/graphs.
 - Can simplify things by re-ifying an irrelevant model.
 - Win for automation.
- Automation pays off when reasoning about separation logic.

Outlook

Future

- Still a fair amount of work for a more realistic system.
 - Reasoning about concurrency.
 - Brookes '07, Appel '08, Nanevski '09
 - Reasoning about failures.
 - Proofs can still be tedious & long.
 - Domain specific external provers.

Code Slides

Implementation: insert

```
get (H h, int index, #list int# m)
  : Cmd (llist h m)
    (r  $\Rightarrow$  llist h m * [r = nth m index])
{
  let hd := *h in
  // Extract the index element from the list from hd to Null
  while (hd != Null) {

    let nde := *hd in
    if (index == 0) return (Some nde.val);
    hd := nde.next;
    index--;
  }
  return None;
}
```


Implementation: insert

```
get (H h, int index, #list int# m)
  : Cmd (llist h m)
    (r  $\Rightarrow$  llist h m * [r = nth m index])
{
  let hd := *h in
  // Extract the index element from the list from hd to Null
  while (hd != Null) {
    // Need to specify the loop invariant
    let nde := *hd in
    if (index == 0) return (Some nde.val);
    hd := nde.next;
    index--;
  }
  return None;
}
```

Implementation: insert

```

get (H h, int index, #list int# m)
  : Cmd (llist h m)
    (r  $\Rightarrow$  llist h m * [r = nth m index])
{
  let hd := *h in
  Fix3 (fun hd j m  $\Rightarrow$  llseg hd m Null)
    (fun hd j m (r : option int)  $\Rightarrow$  llseg hd m Null *
      [r = nth m j])
  (fun self hd j m  $\Rightarrow$ 
    IfNull hd Then
      {{ Return None }}
    Else
      let nde := *hd in
      IfZero j Then
        {{ Return (Some (val nde)) }}
      Else
        {{ self (next nde) j (tail m) <@> _ }})
  hd i m <@> _
}

```